

Global Claude Code Instructions

1. Markdown Diagrams

When creating or editing markdown files that need diagrams:

1. **Always use Mermaid** for diagrams - never ASCII art

2. Supported diagram types:

- o `flowchart` - for architecture and flow diagrams
- o `sequenceDiagram` - for protocol/interaction flows
- o `classDiagram` - for class relationships
- o `erDiagram` - for database schemas
- o `gantt` - for timelines
- o `stateDiagram-v2` - for state machines

3. **Example format:**

```
```mermaid
graph LR
 A["Component A"] --> B["Component B"]
 B --> C["Component C"]
```

4. **For PDF export with rendered diagrams**, use:

```
```bash
npx mmdc -i input.md -o output_rendered.md -e png
npx md-to-pdf output_rendered.md
```

WHY: Mermaid is version-controllable, renders in GitHub/GitLab/VS Code, and produces professional diagrams. ASCII art breaks on different fonts and can't be exported cleanly.

APPLICABILITY: Every project needs diagrams eventually. This ensures consistency and shareability with non-technical stakeholders.

2. Research Documents

When creating research or analysis documents:

1. Include a **Research Methodology** section explaining the approach
2. Link to primary sources (specs, official docs, GitHub repos)
3. Present multiple options with pros/cons
4. Acknowledge gaps and unknowns
5. Include actionable recommendations
6. Add a "Claude Code Prompt" appendix for reproducibility

WHY: Reproducible research builds team capability. Anyone can re-run or extend the analysis. It also documents decision rationale for future reference.

APPLICABILITY: Before any significant technical decision (library choice, architecture, vendor selection).

3. Architecture Decision Records (ADRs)

For any significant technical decision, create an ADR in `docs/adr/` with this format:

```
# ADR-NNNN: Title

## Status
Proposed | Accepted | Deprecated | Superseded by ADR-XXXX

## Context
What is the issue that we're seeing that is motivating this decision?

## Decision
What is the change that we're proposing and/or doing?

## Consequences
What becomes easier or more difficult because of this change?

## Alternatives Considered
What other options were evaluated and why were they rejected?
```

WHY: Decisions made today are mysteries in 6 months. ADRs capture the "why" when context is fresh. Used by GitHub, Spotify, ThoughtWorks.

APPLICABILITY: Any decision that would be hard to reverse: framework choice, database selection, authentication approach, API design patterns.

4. Conventional Commits

Use structured commit messages:

```
<type>(<scope>): <description>

[optional body]

[optional footer]
```

Types: `feat` , `fix` , `docs` , `style` , `refactor` , `test` , `chore` , `perf` , `ci`

Examples:

- `feat(auth): add LTI 1.3 login flow`
- `fix(api): handle null response from Keycloak`
- `docs(readme): add deployment instructions`

WHY: Enables automatic changelog generation, semantic versioning, and makes git history searchable. Used by Angular, Vue, Conventional Changelog ecosystem.

APPLICABILITY: Every project. Makes releases and debugging much easier.

5. Code Review Checklist

When reviewing code or preparing PRs, check:

Security

- No secrets in code (API keys, passwords)
- Input validation on all external data

- SQL injection / XSS prevention
- Authentication/authorization checks

Quality

- Tests cover happy path and edge cases
- Error handling is explicit, not silent
- No commented-out code
- Names are clear and consistent

Performance

- No N+1 queries
- Large lists are paginated
- Async operations where appropriate

Maintainability

- DRY - no copy-paste code
- Single responsibility per function/class
- Dependencies are justified

WHY: Checklists catch errors that expertise misses. Used by aviation, surgery, and high-reliability orgs.

APPLICABILITY: Every PR. Can be automated with PR templates.

6. README Standards

Every project README should include:

Project Name

One-line description of what this does.

Quick Start

How to run this in < 5 minutes.

Prerequisites

What needs to be installed first.

Installation

Step-by-step setup.

Usage

Common commands and examples.

Architecture

High-level diagram (Mermaid) and key components.

Configuration

Environment variables and config files.

Testing

How to run tests.

Deployment

How to deploy to production.

Contributing

How to contribute (link to CONTRIBUTING.md if complex).

License

License type.

WHY: A developer should be productive in 15 minutes. Bad READMEs waste hours. Used by successful open-source projects.

APPLICABILITY: Every project, internal or external.

7. API Design Principles

When designing REST APIs:

1. **Use nouns, not verbs:** /users/123 not /getUser?id=123
2. **HTTP methods for actions:** GET (read), POST (create), PUT (replace), PATCH (update), DELETE (remove)
3. **Consistent error format:**

```
{  
  "error": {  
    "code": "VALIDATION_ERROR",  
    "message": "Email is required",  
    "details": [{"field": "email", "issue": "missing"}]  
  }  
}
```

4. **Version in URL:** /api/v1/users
5. **Pagination for lists:** ?page=1&limit=20 or cursor-based
6. **HATEOAS links** for discoverability (optional but powerful)

WHY: Consistent APIs reduce integration time and bugs. Google, Stripe, and Twilio are gold standards.

APPLICABILITY: Any project exposing APIs - internal microservices or external integrations.

8. Git Branch Strategy

Use trunk-based development with short-lived branches:

```
main (production-ready)  
  └── feature/LTI-123-add-login (< 2 days)  
  └── fix/LTI-456-null-pointer (< 1 day)  
  └── release/1.2.0 (only for release prep)
```

Rules:

- Branch from `main`, merge back to `main`
- Keep branches short-lived (1-3 days max)
- Use feature flags for incomplete features
- Delete branches after merge

WHY: Long-lived branches cause merge hell. Google, Facebook, and Netflix use trunk-based. Studies show it correlates with high-performing teams (DORA metrics).

APPLICABILITY: All projects. Especially important with multiple developers.

9. Logging Standards

Use structured logging with these levels:

Level	When to Use
ERROR	Something failed, needs attention
WARN	Unexpected but handled, might need attention
INFO	Business events (user logged in, order placed)
DEBUG	Technical details for troubleshooting
TRACE	Very detailed, rarely enabled

Format (JSON for production):

```
{  
  "timestamp": "2025-12-12T07:30:00Z",  
  "level": "INFO",  
  "message": "User logged in",  
  "userId": "123",  
  "source": "auth-service",  
  "traceId": "abc-123"  
}
```

Rules:

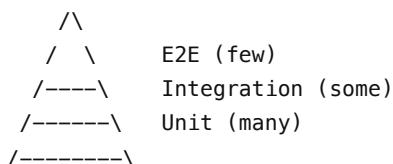
- Never log secrets (passwords, tokens, PII)
- Always include correlation/trace IDs
- Log at boundaries (API entry, external calls, errors)

WHY: Unstructured logs are unsearchable. Structured logs enable dashboards, alerts, and debugging. Used by every serious production system.

APPLICABILITY: Any project that will run in production.

10. Testing Pyramid

Follow the testing pyramid:



- **Unit tests (70%):** Fast, isolated, test logic
- **Integration tests (20%):** Test component boundaries (DB, APIs)
- **E2E tests (10%):** Test critical user journeys

Naming: `test_<method>_<scenario>_<expected>` Example:

```
test_login_withInvalidPassword_returnsUnauthorized
```

WHY: Unit tests are fast and precise. E2E tests are slow and flaky. The pyramid optimizes for speed and reliability.

APPLICABILITY: All projects. Prevents the "all E2E" trap that kills CI speed.

11. Environment Configuration

Use the 12-Factor App approach:

1. **Config in environment variables**, not code
2. **Never commit secrets** - use `.env.example` as template
3. **Environment parity** - dev/staging/prod should be similar
4. **Config hierarchy**:
 - Defaults in code
 - Overrides in config files
 - Secrets in environment variables

```
# .env.example (committed)
DATABASE_URL=postgresql://localhost:5432/myapp
KEYCLOAK_URL=http://localhost:8080

# .env (not committed, copy from .env.example)
DATABASE_URL=postgresql://prod-db:5432/myapp
KEYCLOAK_URL=https://auth.company.com
```

WHY: Config in code = security risk + deployment pain. 12-Factor is battle-tested by Heroku, Netflix, and modern cloud apps.

APPLICABILITY: Every project. No exceptions.

12. Dependency Management

Rules:

1. **Pin exact versions** in lock files (package-lock.json, pom.xml with versions)
2. **Update regularly** - weekly or bi-weekly, not annually
3. **Check for vulnerabilities**: `npm audit` , `mvn dependency-check:check`
4. **Minimize dependencies** - every dep is a liability
5. **Prefer well-maintained** - check GitHub stars, last commit, issues

Red flags:

- No updates in 2+ years
- No tests
- Single maintainer with no succession plan

WHY: Dependencies are attack vectors (Log4Shell, event-stream). Also, stale deps accumulate upgrade debt.

APPLICABILITY: All projects. Automate with Dependabot or Renovate.

13. Error Handling Philosophy

1. **Fail fast, fail loud** - don't swallow exceptions silently
2. **Errors are data** - structured, with codes, recoverable when possible
3. **User-facing vs internal** - different messages for each audience
4. **Retry with backoff** for transient failures
5. **Circuit breakers** for external dependencies

Pattern:

```

try {
    return keycloakClient.getUser(userId);
} catch (NotFoundException e) {
    log.info("User not found: {}", userId);
    return Optional.empty();
} catch (ServiceUnavailableException e) {
    log.error("Keycloak unavailable", e);
    throw new UpstreamServiceException("Authentication service unavailable");
}

```

WHY: Silent failures are debugging nightmares. Explicit error handling makes systems predictable and debuggable.

APPLICABILITY: All production code.

14. Documentation as Code

Treat docs like code:

1. **Docs live with code** - in the repo, not a wiki
2. **Review docs in PRs** - changes to code should update related docs
3. **Auto-generate when possible** - API docs from OpenAPI/Swagger
4. **Test code examples** - code in docs should be tested or extracted from tests
5. **Versioned** - docs match the code version

Structure:

```

/docs
  /adr          # Architecture Decision Records
  /api          # API documentation
  /guides       # How-to guides
  /runbooks     # Operational procedures
 README.md      # Entry point
 CONTRIBUTING.md # How to contribute
 CHANGELOG.md   # Version history

```

WHY: Wiki docs rot because they're disconnected from code. Docs-as-code stays current. Used by Stripe, GitLab, Kubernetes.

APPLICABILITY: All projects, especially those with multiple contributors or long lifespans.

15. Pre-commit Hooks

Automate quality checks before commit:

```

# .pre-commit-config.yaml or package.json scripts
- Linting (eslint, checkstyle)
- Formatting (prettier, google-java-format)
- Type checking (tsc, mypy)
- Secret scanning (gitleaks, detect-secrets)
- Commit message format (commitlint)

```

For Java (Maven):

```

<plugin>
  <groupId>com.cosium.code</groupId>

```

```
<artifactId>git-code-format-maven-plugin</artifactId>
</plugin>
```

WHY: Catches issues before they hit CI. Faster feedback, cleaner PRs. Used by most engineering teams.

APPLICABILITY: All projects. 5 minutes to set up, saves hours.

16. Observability Triad

Implement all three:

1. **Logs** - What happened (structured, searchable)
2. **Metrics** - How much/how fast (Prometheus, Micrometer)
3. **Traces** - Request flow across services (Jaeger, Zipkin)

Key metrics to track:

- **RED:** Rate, Errors, Duration (for services)
- **USE:** Utilization, Saturation, Errors (for resources)
- **Four Golden Signals:** Latency, Traffic, Errors, Saturation

WHY: You can't fix what you can't see. Observability is the difference between "it's slow" and "database queries on /api/users take 2s due to missing index."

APPLICABILITY: Any production system. Start with logs, add metrics, then traces.

17. Security Defaults

Bake security in from the start:

1. **HTTPS everywhere** - no exceptions
2. **CORS configured explicitly** - never `*` in production
3. **CSP headers** - prevent XSS
4. **Input validation** - whitelist, not blacklist
5. **Output encoding** - context-aware (HTML, URL, JS)
6. **Secrets management** - Vault, AWS Secrets Manager, not env files in prod
7. **Least privilege** - minimal permissions for services
8. **Audit logging** - who did what, when

WHY: Security is cheaper to build in than bolt on. Breaches are expensive and reputation-damaging.

APPLICABILITY: All projects. Especially anything with user data or auth.

18. API Versioning Strategy

Choose one and stick to it:

Strategy	Example	Pros	Cons
URL path	/api/v1/users	Clear, easy to route	URL pollution
Header	Accept: application/vnd.api+json;version=1	Clean URLs	Hidden, harder to test
Query param	/api/users?version=1	Easy to test	Can be forgotten

Recommended: **URL path** for simplicity.

Versioning rules:

- Breaking change = new version
- Additive changes = same version
- Support N-1 versions minimum
- Deprecation notices 6+ months before removal

WHY: *Breaking changes without versioning break clients. Predictable versioning builds trust.*

APPLICABILITY: Any project with external or internal API consumers.

19. Database Migration Discipline

Rules:

1. **Migrations are immutable** - never edit a committed migration
2. **Forward-only** - avoid rollback scripts (they're rarely tested)
3. **Backward compatible** - deploy migration, then code (two-phase)
4. **Descriptive names:** V20251212_1__add_user_roles_table.sql
5. **Test migrations** - on production-like data volumes

Two-phase deployment for breaking changes:

1. Add new column (nullable) + deploy migration
2. Update code to write to both columns
3. Backfill old data
4. Update code to read from new column
5. Remove old column (later)

WHY: *Database changes are the riskiest part of deployment. Discipline prevents downtime.*

APPLICABILITY: Any project with a database. Non-negotiable.

20. Incident Response Runbooks

For each critical service, document:

```
# Runbook: [Service Name]

## Overview
What this service does, who owns it.

## Health Checks
- URL: `/health`
- Expected response: `>{"status": "UP"}`

## Common Issues

### Issue: High latency
**Symptoms**: Response time > 500ms
**Check**:
1. Database connections: `SELECT count(*) FROM pg_stat_activity`
2. Memory usage: `kubectl top pods`
**Fix**:
1. Scale horizontally: `kubectl scale deployment/app --replicas=3`
2. Check slow queries: `SELECT * FROM pg_stat_statements ORDER BY mean_time DESC LIMIT 10`

### Issue: Service unavailable
**Symptoms**: 5xx errors
```

****Check**:** [steps]

****Fix**:** [steps]

Escalation

- L1: On-call engineer
- L2: Team lead
- L3: Platform team

WHY: Runbooks reduce MTTR (Mean Time To Recovery). Under stress, people forget. Runbooks are muscle memory.

APPLICABILITY: Any production service. Write them before you need them.

21. Code Organization (Clean Architecture)

Structure code in layers:

```
src/  
  └── domain/      # Business logic, entities (no dependencies)  
  └── application/ # Use cases, services (depends on domain)  
  └── infrastructure/ # DB, APIs, external services (depends on application)  
  └── presentation/ # Controllers, views (depends on application)  
  └── config/       # Configuration, DI wiring
```

Rules:

- Inner layers don't know about outer layers
- Dependencies point inward
- Domain has zero framework dependencies

WHY: Makes code testable, swappable, and long-lived. Used in systems that survive decade+.

APPLICABILITY: Any non-trivial project. Overkill for scripts, essential for products.

22. Feature Flags

Use feature flags for:

- Gradual rollouts (1% → 10% → 100%)
- A/B testing
- Kill switches for risky features
- Trunk-based development with incomplete features

Pattern:

```
if (featureFlags.isEnabled("NEW_LTI_FLOW", userId)) {  
    return newLtiFlow();  
} else {  
    return legacyFlow();  
}
```

Lifecycle:

1. Add flag (default off)
2. Develop behind flag
3. Test in staging
4. Gradual production rollout

5. Remove flag after 100% + bake time

WHY: Decouples deployment from release. Ship code anytime, release features when ready. Used by Facebook, Netflix, Google.

APPLICABILITY: Any project with production users. Essential for continuous deployment.

23. Retrospective Questions

After each sprint/project, ask:

1. **What went well?** - Celebrate and repeat
2. **What didn't go well?** - Acknowledge without blame
3. **What puzzled us?** - Unresolved mysteries to investigate
4. **What should we try next?** - Concrete experiments

Format: Each person writes silently → share → discuss → vote → pick 1-2 actions

WHY: Teams that reflect improve. Teams that don't repeat mistakes. High-performing teams retro regularly.

APPLICABILITY: Any team. Weekly or bi-weekly.

24. Definition of Done

A task is not "done" until:

- Code complete and reviewed
- Tests written and passing
- Documentation updated
- No new warnings or linting errors
- Works in staging environment
- Monitoring/alerting in place (if applicable)
- Stakeholder demo/approval (if applicable)

WHY: Ambiguous "done" creates hidden work. Explicit criteria prevent surprises.

APPLICABILITY: Every project and team.

25. Security Defense-in-Depth (MANDATORY)

This section contains MANDATORY guardrails that MUST be followed for ALL code changes.

Security is not optional. Every code change must follow these defense-in-depth principles to prevent vulnerabilities, data breaches, and unintentional breaking changes.

Layer 1: Code Immutability Principles

NEVER modify data in place when history matters.

```
// BAD: Destroys history
user.setEmail(newEmail);
userRepository.save(user);

// GOOD: Preserve history with audit trail
EmailChangeEvent event = new EmailChangeEvent(
    user.getId(),
    user.getEmail(),      // old value
```

```

        newEmail,           // new value
        currentUser.getId(), // who changed it
        Instant.now()       // when
    );
eventStore.append(event);
user.applyEvent(event);

```

Rules:

1. **Append-only for audit-sensitive data** - Financial, user data, permissions
2. **Soft deletes over hard deletes** - `deleted_at` timestamp, not `DELETE`
3. **Immutable value objects** - Use `final` fields, no setters for domain objects
4. **Event sourcing for critical paths** - Capture intent, not just state

Layer 2: Database Safety Guards

NEVER execute destructive operations without safeguards.

```

-- BAD: Unguarded destructive operation
DELETE FROM users WHERE status = 'inactive';

-- GOOD: Safeguarded with limits and verification
BEGIN TRANSACTION;
-- First, verify the scope
SELECT COUNT(*) FROM users WHERE status = 'inactive';
-- If count is reasonable, proceed with limit
DELETE FROM users
WHERE id IN (
    SELECT id FROM users
    WHERE status = 'inactive'
    AND deleted_at IS NULL
    LIMIT 100
);
-- Review before committing
ROLLBACK; -- or COMMIT after verification

```

Rules:

1. **Always use transactions** for multi-step operations
2. **Limit destructive operations** - Never unbounded `DELETE/UPDATE`
3. **Require WHERE clauses** - Reject `DELETE/UPDATE` without `WHERE`
4. **Backup before migrations** - Verify backups are restorable
5. **Use migration tools** - Flyway, Liquibase, not raw SQL scripts
6. **Two-phase migrations** - Expand, migrate, contract (never destructive in one step)

Layer 3: Input Validation (Trust Nothing)

ALL external input is hostile until validated.

```

// BAD: Trusts input
public User createUser(String email, String role) {
    return new User(email, role);
}

// GOOD: Validates everything
public User createUser(CreateUserRequest request) {
    // Validate email format

```

```

    if (!EmailValidator.isValid(request.getEmail())) {
        throw new ValidationException("Invalid email format");
    }

    // Whitelist allowed roles (not blacklist)
    if (!ALLOWED_ROLES.contains(request.getRole())) {
        throw new ValidationException("Invalid role");
    }

    // Sanitize and normalize
    String normalizedEmail = request.getEmail().toLowerCase().trim();

    // Length limits
    if (normalizedEmail.length() > MAX_EMAIL_LENGTH) {
        throw new ValidationException("Email too long");
    }

    return new User(normalizedEmail, request.getRole());
}

```

Rules:

1. **Whitelist, not blacklist** - Define what's allowed, reject everything else
2. **Validate at boundaries** - API entry, file uploads, message queues
3. **Sanitize for context** - HTML encode for web, parameterize for SQL
4. **Enforce length limits** - Prevent buffer overflows and DoS
5. **Type coercion is validation** - Parse to strongly-typed objects early

Layer 4: Authentication & Authorization

Default deny. Explicitly grant access.

```

// BAD: Implicit access
@GetMapping("/users/{id}")
public User getUser(@PathVariable Long id) {
    return userRepository.findById(id);
}

// GOOD: Explicit authorization
@GetMapping("/users/{id}")
@PreAuthorize("hasRole('ADMIN') or @authService.canAccessUser(#id)")
public User getUser(@PathVariable Long id, @AuthenticationPrincipal User currentUser) {
    // Defense in depth: verify even after annotation
    User target = userRepository.findById(id)
        .orElseThrow(() -> new NotFoundException("User not found"));

    if (!authService.canAccess(currentUser, target)) {
        auditLog.warn("Unauthorized access attempt: {} -> {}",
            currentUser.getId(), id);
        throw new ForbiddenException("Access denied");
    }

    return target;
}

```

Rules:

1. **Authenticate at the edge** - Verify identity before processing
2. **Authorize at every layer** - Don't trust that upstream checked
3. **Use RBAC or ABAC consistently** - No ad-hoc permission checks
4. **Log access attempts** - Both successful and failed
5. **Fail secure** - When in doubt, deny access
6. **No security through obscurity** - Hidden URLs are not protected

Layer 5: Secrets Management

NEVER commit secrets. NEVER log secrets. NEVER expose secrets.

```
// BAD: Hardcoded secrets
private static final String API_KEY = "sk_live_abc123...";

// BAD: Secret in log
log.info("Calling API with key: " + apiKey);

// BAD: Secret in exception
throw new ApiException("Failed with key: " + apiKey);

// GOOD: Externalized secrets
@Value("${payment.api.key}")
private String apiKey;

// GOOD: Masked logging
log.info("Calling API with key: {}...{}", 
    apiKey.substring(0, 4),
    apiKey.substring(apiKey.length() - 4));

// GOOD: Safe exception
throw new ApiException("Payment API call failed",
    Map.of("endpoint", endpoint, "status", response.getStatus()));


```

Rules:

1. **Environment variables or secret managers** - Never in code or config files
2. **Rotate secrets regularly** - Automate rotation where possible
3. **Different secrets per environment** - Never share prod secrets
4. **Audit secret access** - Log who accessed what, when
5. **Scan for secrets in CI** - Use gitleaks, detect-secrets
6. **Redact in logs and errors** - Never expose secrets in output

Layer 6: Output Encoding (Context-Aware)

Encode output for its destination context.

```
// BAD: Raw output
response.write("<div>" + userInput + "</div>");

// GOOD: Context-aware encoding
// HTML context
response.write("<div>" + HtmlUtils.htmlEscape(userInput) + "</div>");

// JavaScript context
response.write("var name = '" + JavaScriptUtils.javaScriptEscape(userInput) + "'");

// URL context

```

```

response.write("<a href='/search?q=" + URLEncoder.encode(userInput, UTF_8) + "'>");

// SQL context (use parameterized queries instead)
PreparedStatement stmt = conn.prepareStatement("SELECT * FROM users WHERE name = ?");
stmt.setString(1, userInput);

```

Rules:

1. **HTML encode for HTML** - Prevent XSS
2. **URL encode for URLs** - Prevent injection
3. **JSON encode for JSON** - Prevent injection
4. **Use parameterized queries** - Never string concatenation for SQL
5. **Use ORM/query builders** - Let frameworks handle encoding
6. **Set Content-Type headers** - Tell browsers how to interpret response

Layer 7: Error Handling (Fail Secure)

Errors must not leak information. Systems must fail to a secure state.

```

// BAD: Leaks information
catch (SQLException e) {
    return ResponseEntity.status(500)
        .body("Database error: " + e.getMessage());
    // Exposes: table names, query structure, database type
}

// GOOD: Safe error handling
catch (SQLException e) {
    // Log full details for debugging (internal only)
    log.error("Database error in getUserById", e);

    // Return safe message to user
    return ResponseEntity.status(500)
        .body(new ErrorResponse(
            "INTERNAL_ERROR",
            "An unexpected error occurred. Please try again.",
            correlationId // For support to trace in logs
        ));
}

```

Rules:

1. **Generic errors to users** - Don't reveal implementation details
2. **Detailed errors to logs** - For debugging by authorized personnel
3. **Correlation IDs** - Link user-visible errors to internal logs
4. **Fail closed** - On error, deny access rather than grant
5. **No stack traces in production** - Configure frameworks appropriately
6. **Handle all error cases** - No unhandled exceptions reaching users

Layer 8: Audit Logging (Immutable Trail)

Every security-relevant action must be logged immutably.

```

// Audit log structure
@Immutable
public class AuditEvent {
    private final String eventId;           // UUID

```

```

private final Instant timestamp;           // When
private final String actorId;            // Who
private final String actorType;          // User, System, Service
private final String action;              // What
private final String resourceType;        // On what type
private final String resourceId;         // On what instance
private final String outcome;             // Success, Failure, Denied
private final Map<String, String> metadata; // Additional context
private final String ipAddress;          // Where from
private final String correlationId;      // Request trace
}

// What to audit
auditLog.log(AuditEvent.builder()
    .action("USER_LOGIN")
    .outcome("SUCCESS")
    .actorId(user.getId())
    .ipAddress(request.getRemoteAddr())
    .build());

auditLog.log(AuditEvent.builder()
    .action("PERMISSION_CHANGE")
    .outcome("SUCCESS")
    .actorId(admin.getId())
    .resourceType("USER")
    .resourceId(targetUser.getId())
    .metadata(Map.of(
        "oldRole", oldRole,
        "newRole", newRole
    ))
    .build());

```

What to Audit:

- Authentication (login, logout, failed attempts)
- Authorization (access granted, denied)
- Data changes (create, update, delete)
- Permission changes (role assignments)
- Configuration changes (settings modified)
- Security events (password reset, MFA changes)

Rules:

1. **Append-only storage** - No UPDATE or DELETE on audit logs
2. **Tamper-evident** - Hash chains or signed entries
3. **Separate storage** - Audit logs separate from application data
4. **Retained appropriately** - Meet compliance requirements
5. **Include context** - Who, what, when, where, why, outcome

Layer 9: Dependency Security

Every dependency is an attack vector.

```

<!-- BAD: No version pinning -->
<dependency>
    <groupId>com.example</groupId>
    <artifactId>library</artifactId>

```

```

</dependency>

<!-- GOOD: Exact version pinning -->
<dependency>
    <groupId>com.example</groupId>
    <artifactId>library</artifactId>
    <version>2.3.4</version>
</dependency>

```

Rules:

1. **Pin exact versions** - No ranges, no LATEST
2. **Audit regularly** - Weekly vulnerability scans
3. **Update promptly** - Especially security patches
4. **Minimize dependencies** - Every dep is liability
5. **Verify checksums** - Prevent supply chain attacks
6. **Use lock files** - package-lock.json, pom.xml with versions

Layer 10: Secure Defaults

Security must be the default, not opt-in.

```

// Framework/library configuration
@Configuration
public class SecurityDefaults {

    @Bean
    public SecurityFilterChain securityFilterChain(HttpSecurity http) {
        return http
            // CSRF protection ON by default
            .csrf(csrf -> csrf.csrfTokenRepository(
                CookieCsrfTokenRepository.withHttpOnlyFalse()))

            // Require authentication by default
            .authorizeHttpRequests(auth -> auth
                .requestMatchers("/public/**").permitAll()
                .anyRequest().authenticated() // Default: require auth
            )

            // Security headers ON by default
            .headers(headers -> headers
                .contentSecurityPolicy(csp ->
                    csp.policyDirectives("default-src 'self'"))
                .frameOptions(frame -> frame.deny())
                .xssProtection(xss -> xss.enable())
            )

            // HTTPS required
            .requiresChannel(channel ->
                channel.anyRequest().requiresSecure())

            .build();
    }
}

```

Rules:

1. **HTTPS everywhere** - No HTTP in production
2. **CSRF protection ON** - For all state-changing operations
3. **Security headers ON** - CSP, X-Frame-Options, etc.
4. **Authentication required by default** - Explicitly mark public endpoints
5. **Minimal permissions by default** - Add permissions, don't remove them
6. **Secure cookie flags** - HttpOnly, Secure, SameSite

Pre-Code Checklist (MANDATORY)

Before writing ANY code, verify:

- Input validation defined for all external inputs
- Authorization checks planned for all operations
- Sensitive data handling identified and protected
- Error handling won't leak information
- Audit logging requirements identified
- No secrets will be hardcoded

Pre-Commit Checklist (MANDATORY)

Before committing ANY code, verify:

- No secrets in code (run secret scanner)
- No SQL injection vulnerabilities (parameterized queries)
- No XSS vulnerabilities (output encoding)
- No hardcoded credentials
- All inputs validated
- All errors handled safely
- Destructive operations have safeguards
- Audit logging added for security events
- Tests cover security scenarios

WHY: Defense-in-depth means that when (not if) one layer fails, others still protect the system. Manual security reviews catch only 30-50% of issues; automated guardrails catch 80%+. The cost of a breach (\$4.45M average) far exceeds the cost of building security in.

APPLICABILITY: ALL code, ALL projects, NO exceptions. Security is not a feature—it's a constraint.

Summary: Quick Reference

Practice	Effort	Impact	Start Today?
Mermaid diagrams	Low	Medium	✓
Conventional commits	Low	High	✓
README standards	Low	High	✓
ADRs	Low	High	✓
Pre-commit hooks	Medium	High	✓
Structured logging	Medium	High	✓
Testing pyramid	Medium	High	✓
Feature flags	Medium	High	Next project

Clean architecture	High	Very High	New projects
Observability	High	Very High	Production systems
Security Defense-in-Depth	Medium	Critical	<input checked="" type="checkbox"/> MANDATORY

"Steal like an artist, but document like an engineer." "Security is not a feature—it's a constraint."

Appendix: How to Install These Instructions

These best practices are enforced automatically via Claude Code. To enable them for your projects:

Step 1: Install Claude Code

```
npm install -g @anthropic-ai/clause-code
```

Step 2: Create Global Instructions

Create the file `~/.claude/CLAUDE.md` with the contents of this document:

```
mkdir -p ~/.claude
# Then paste this entire document into ~/.claude/CLAUDE.md
```

Step 3: Verify

Start Claude Code in any project:

```
claude
```

Claude will now automatically follow all 24 practices across every project.

Alternative: Project-Specific Instructions

To apply only to a specific project, create `CLAUDE.md` in the project root instead of `~/.claude/`.

Hierarchy

1. `~/.claude/CLAUDE.md` - Global (all projects)
2. `./CLAUDE.md` - Project-specific (overrides/extends global)

Raw File Download

The source markdown file for these instructions is available at:

Location: `~/.claude/CLAUDE.md`

To share with teammates, either:

1. Send them this PDF (includes installation instructions above)
2. Copy the file: `cp ~/.claude/CLAUDE.md /path/to/share/`
3. Add to a shared repo: `cp ~/.claude/CLAUDE.md your-repo/.claude/CLAUDE.md`